

SYNOPSIS OF EXPLORATORY PROJECT (EP)

ON

ADSLM – ADAPTIVE DATA SCIENCE

LANGUAGE MODEL

submitted in partial fulfilment of the requirements for the award of degree of

BACHELOR OF ENGINEERING

In

COMPUTER SCIENCE AND ENGINEERING

Submitted by:

KARTIKEYA JAIN

2310991865

Supervised By:

MONIKA AGGARWAL

ASSISTANT PROFESSOR



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

CHITKARA UNIVERSITY INSTITUTE OF ENGINEERING AND TECHNOLOGY

CHITKARA UNIVERSITY, PUNJAB, INDIA

CONTENTS

S.NO	TITLE	PAGE
	Abstract	1
1	Introduction	2
1.1	Background	2
1.2	Problem Statement	3
1.3	Need of the Project	3
1.4	Objectives	4
1.5	Scope of the Project	4
2	Methodology	5
2.1	System Architecture	5
2.2	Dataset Analysis	6
2.3	Machine Learning Task Detection	6
2.4	Data Preprocessing	7
2.5	Model Recommendation and Training	7
2.6	Model Evaluation	8
2.7	Explainable AI	8
2.8	Adaptive Insight Generation	9
2.9	Report Generation	9
3	Tools and Technologies	10
3.1	Software Architecture	10
3.2	Sample Dataset and Industrial Use Case	11
3.3	Frontend and Backend	11
4	Project Plan	12
4.1	Expected Outcomes	13
4.2	Future Scope	13
	Conclusion	14

ABSTRACT

The rapid growth of industrial data generated through sensors, machines, maintenance systems, and production processes has created significant opportunities for data-driven decision-making. However, transforming raw industrial data into meaningful machine learning results generally requires several manual steps, including data analysis, preprocessing, task identification, model selection, training, evaluation, interpretation, and reporting. These activities require considerable technical expertise and time.

ADSLM (Adaptive Data Science Language Model) is proposed as an intelligent AutoML and Data Science Copilot designed to automate this end-to-end workflow. The system accepts a CSV dataset as input and automatically analyzes its structure, identifies data quality issues, detects the appropriate machine learning task, performs preprocessing, recommends suitable machine learning algorithms, trains and compares multiple models, evaluates their performance, and selects the best-performing model.

The system further incorporates an Explainable AI (XAI) component to identify important features influencing model predictions. Instead of presenting only technical metrics, ADSLM generates natural-language insights and adapts these explanations according to three expertise levels: Beginner, Intermediate, and Expert. This makes machine learning results understandable to both technical and non-technical stakeholders.

The proposed system uses Python, Pandas, NumPy, Scikit-learn, XGBoost, FastAPI, Streamlit, Plotly, Joblib, and ReportLab. A predictive maintenance dataset containing industrial sensor parameters such as temperature, vibration, pressure, RPM, oil level, and maintenance information is used as the primary demonstration use case. The target variable represents machine fault conditions, making the problem suitable for classification.

The major contribution of ADSLM is its ability to adapt at three levels: dataset, machine learning task, and user expertise. The system provides a complete workflow from raw data to model, explanation, insights, and report, making it suitable for industrial data science applications and future integration with large-scale production environments.

1. INTRODUCTION

1.1 Background

Modern industries generate large volumes of data from sensors, machines, maintenance operations, energy systems, production processes, and quality-control mechanisms. Machine learning can transform this data into useful predictions and insights, such as identifying equipment faults, detecting anomalies, forecasting energy consumption, and predicting maintenance requirements.

However, developing a machine learning solution traditionally involves a sequence of technical activities. A data scientist must first understand the dataset, identify its characteristics, clean missing and inconsistent values, select relevant features, determine the type of machine learning problem, select appropriate algorithms, train multiple models, evaluate their performance, interpret the results, and communicate the findings to stakeholders.

ADSLM addresses this challenge by integrating these activities into a single automated pipeline. The system is designed to transform a raw CSV dataset into a complete machine learning analysis consisting of dataset profiling, task detection, preprocessing, model recommendation, training, evaluation, explainability, natural-language insights, and report generation.

The system is particularly motivated by industrial applications such as predictive maintenance, where machine sensor readings can be used to identify different fault conditions and support proactive maintenance decisions.

1.2 Problem Statement

The conventional machine learning development process is highly dependent on manual intervention and specialized technical knowledge. Different datasets may have different structures, data types, missing values, outliers, categorical variables, and prediction objectives. Consequently, a fixed machine learning workflow may not be suitable for every dataset.

1. Manual analysis of newly received datasets.
2. Difficulty in automatically identifying the appropriate ML task.
3. Time-consuming data preprocessing.
4. Manual selection and comparison of machine learning models.
5. Difficulty in interpreting model predictions.
6. Limited accessibility of ML results to non-technical users.
7. Manual preparation of reports and documentation.
8. Lack of a unified workflow connecting all these stages.

ADSLM aims to overcome these limitations by providing an automated and adaptive data science pipeline.

1.3 Need of the Project

Industrial organizations work with datasets that may contain missing values, duplicates, outliers, mixed data types, categorical variables, and complex relationships between features. Manually preparing such datasets for machine learning can be time-consuming and error-prone.

- Reduce repetitive data science tasks.
- Automate dataset profiling and preprocessing.
- Identify suitable machine learning problem types.
- Compare multiple machine learning models.
- Provide meaningful evaluation metrics.
- Explain important features influencing predictions.
- Present results in an understandable manner.
- Generate structured reports automatically.
- Support both technical and non-technical users.

In ABB-like industrial systems, such a platform can support applications including machine fault prediction, anomaly detection, energy management, and quality-related classification.

1.4 Objectives

The primary objective of the project is to develop an intelligent and adaptive data science assistant capable of automating the complete machine learning workflow for structured CSV datasets.

1. To develop an automated dataset analysis module.
2. To identify data types, missing values, duplicates, and outliers.
3. To automatically detect the appropriate machine learning task.
4. To automate data preprocessing.
5. To recommend suitable machine learning algorithms based on dataset characteristics.
6. To train and compare multiple machine learning models.
7. To evaluate models using task-specific performance metrics.
8. To automatically select and save the best-performing model.
9. To provide feature-importance-based explainability.
10. To generate natural-language insights from machine learning results.
11. To adapt explanations according to user expertise.
12. To generate structured TXT, JSON, and PDF reports.
13. To provide an interactive web-based dashboard.
14. To develop the system using a modular and extensible architecture.

1.5 Scope of the Project

The scope of ADSLM includes the development of an end-to-end automated machine learning workflow for structured CSV datasets.

- Dataset profiling and quality analysis.
- Missing-value and duplicate analysis.
- Outlier identification.
- Numerical, categorical, and datetime feature identification.
- Automatic ML task detection.
- Classification, regression, and clustering support.
- Time-series task detection and extensibility.
- Multiple model training and comparison.
- Task-specific performance evaluation.
- Feature-importance-based explainability.

- Natural-language insight generation.
- Expertise-based explanation adaptation.
- Model persistence and report generation.
- Interactive Streamlit visualization.
- FastAPI-based API access.

2. METHODOLOGY

ADSLM follows a modular pipeline architecture, where every stage performs a specific responsibility and passes its output to the next stage. The overall workflow is:

**CSV Upload → Dataset Analysis → Task Detection → Preprocessing → Model Recommendation
→ Training → Evaluation → Best Model Selection → Explainability → Insight Generation →
Expertise Adaptation → Report Generation → Dashboard**

2.1 System Architecture

The system consists primarily of two user-facing components: the FastAPI Backend and the Streamlit Frontend.

FastAPI Backend: The backend acts as the core processing layer. It receives uploaded datasets, executes the machine learning pipeline, coordinates individual services, returns structured JSON responses, saves trained models, generates reports, and provides API documentation.

Streamlit Frontend: The frontend provides an interactive interface through which users can upload CSV datasets, select a target column, select their expertise level, execute the ADSLM pipeline, view dataset information, inspect model performance, analyze feature importance, read generated insights, and download reports.

The backend performs the computational work, while the frontend provides an accessible visualization and interaction layer.

2.2 Dataset Analysis

The first stage of the pipeline is performed by the Dataset Analyzer. After receiving a CSV file, the system analyzes its structure and generates a dataset profile.

- Number of rows and columns.
- Column names and data types.
- Numerical, categorical, and datetime columns.
- Missing values and their percentages.
- Duplicate records and memory usage.
- Summary statistics.
- Outlier detection using the IQR method.
- Possible target-column identification.

For example, a column named Fault, Status, Label, or Target may be identified as a potential target column. This stage provides the foundation for the remaining machine learning workflow.

2.3 Machine Learning Task Detection

After dataset analysis, ADSLM determines the type of machine learning problem using rule-based logic.

Dataset Condition	Detected Task
No target column selected	Clustering
Datetime column + numerical target	Time-Series
Categorical target	Classification
Numerical target	Regression
Target cannot be determined	Unknown

For the predictive-maintenance dataset, the Fault column contains discrete fault categories. Therefore, the system identifies the problem as a classification task.

2.4 Data Preprocessing

Raw industrial datasets may contain missing values, duplicate observations, categorical features, and numerical variables with different scales. Therefore, preprocessing is necessary before model training.

1. Duplicate removal to improve data quality.
2. Feature-target separation.
3. Median imputation for numerical features.
4. Mode imputation for categorical features.
5. One-Hot Encoding for categorical input features.
6. Label Encoding for categorical classification targets.
7. StandardScaler-based feature standardization.
8. 80% training and 20% testing split, with stratification for classification.

Technical consideration: In the current implementation, scaling occurs before the train-test split. For a production-ready implementation, the scaler should be fitted only on the training data and then applied to the test data to prevent potential data leakage. This is identified as a future improvement.

2.5 Model Recommendation and Training

The Model Recommendation Engine recommends machine learning models based on task type, dataset size, feature composition, presence of categorical variables, and dataset characteristics.

Task	Models
Classification	Logistic Regression, Random Forest, SVM, XGBoost
Regression	Linear Regression, Random Forest Regressor, XGBoost Regressor
Clustering	KMeans, DBSCAN
Time-Series	ARIMA and Prophet planned/extensible models

The Training Engine trains the recommended models available in the internal model registry and compares their performance.

Task	Primary Selection Metric	Objective
Classification	F1-Score	Maximize
Regression	RMSE	Minimize
Clustering	Silhouette Score	Maximize
Time-Series	RMSE	Minimize

2.6 Model Evaluation

ADSLM uses task-specific evaluation metrics rather than relying on a single metric for all machine learning problems.

Task	Evaluation Metrics
Classification	Accuracy, Precision, Recall, F1-Score, Confusion Matrix
Regression	RMSE, MAE, R^2
Clustering	Silhouette Score, number of clusters, noise points
Time-Series	RMSE, MAE

F1-score is particularly relevant for fault prediction because it balances precision and recall and can be useful when classes are imbalanced. RMSE and MAE quantify numerical prediction error, R^2 measures explained variance, and Silhouette Score evaluates cluster separation quality.

2.7 Explainable AI

A major feature of ADSLM is that it does not stop at producing a prediction or model score. It attempts to explain model behavior through feature importance.

- `feature_importances_` from tree-based models such as Random Forest and XGBoost.
- Absolute coefficient values from linear models such as Logistic Regression and Linear Regression.
- Ranked top-N important features.
- Model type, strengths, and limitations.

For example, in a predictive maintenance scenario, if vibration and temperature receive high importance scores, the system can identify them as important contributors to the model's fault predictions.

2.8 Adaptive Insight Generation

ADSLM converts technical machine learning results into natural-language explanations through the Insight Generator. Generated insights cover dataset health, detected task, preprocessing operations, model selection, model performance, important features, and actionable recommendations.

Instead of displaying only a technical label such as Classification, the system can explain that the selected target contains discrete categories and that the model is therefore assigning each record to one of those categories.

2.9 Report Generation

After completion of the pipeline, ADSLM generates a structured report using the Report Generator.

- Text report.
- JSON report.
- PDF report when ReportLab is available.

The reports contain the pipeline overview, dataset analysis, task detection, preprocessing summary, model recommendations, model performance, feature importance, generated insights, and actionable recommendations.

3. TOOLS AND TECHNOLOGIES

Technology	Purpose
Python	Primary programming language
Pandas	Dataset loading and manipulation
NumPy	Numerical computation
Scikit-learn	Machine learning algorithms, preprocessing and evaluation
XGBoost	Gradient boosting-based machine learning
FastAPI	Backend API and pipeline orchestration
Streamlit	Interactive frontend dashboard
Plotly	Data visualization
Joblib	Saving and loading trained machine learning models
ReportLab	PDF report generation
JSON	Structured result and report representation
Uvicorn	FastAPI application server

3.1 Software Architecture

The project follows a modular service-based architecture. Each major module is responsible for one stage of the pipeline.

Module	Responsibility
Dataset Analyzer	Dataset profiling and quality analysis
Task Detector	Machine learning task identification
Preprocessing Engine	Data cleaning and transformation
Model Recommendation Engine	Model selection and ranking
Training Engine	Model training and selection
Evaluation Engine	Performance measurement
Explainability Engine	Feature importance and model explanation
Insight Generator	Natural-language insights
Expertise Adapter	User-level explanation adaptation
Report Generator	TXT, JSON and PDF report generation

This modular architecture assigns a specific responsibility to each service, making the system easier to maintain, test, extend, and understand.

3.2 Sample Dataset and Industrial Use Case

The primary demonstration use case is industrial predictive maintenance. The dataset represents machine sensor information and includes parameters such as:

Feature	Description
Machine_ID	Machine identifier
Temperature_C	Machine temperature
Vibration_mm_s	Machine vibration level
Pressure_bar	Pressure reading
RPM	Rotational speed
Oil_Level	Oil level
Maintenance_Hours	Maintenance-related operating information
Fault	Machine fault condition

The Fault attribute acts as the target variable and contains different machine health/fault categories. Since the target is categorical, the problem is treated as a classification problem.

The industrial objective is to learn relationships between machine sensor readings and fault conditions, thereby supporting predictive maintenance and helping identify potentially degraded or faulty equipment

.

3.3 Frontend and Backend

The Streamlit dashboard displays dataset preview, number of rows and columns, missing-value percentage, detected ML task, best model, best score, model recommendations, performance metrics, confusion matrix, model comparison, feature importance, AI-generated insights, actionable recommendations, and report download functionality.

FastAPI provides endpoints for health checking, dataset analysis, and complete pipeline execution, along with automatically generated API documentation.

4. PROJECT PLAN

Phase	Activities	Expected Outcome
Phase 1: Requirement Analysis	Identify industrial problem, objectives, functional requirements and expected outputs	Project requirements finalized
Phase 2: Dataset & Problem Analysis	Study industrial dataset, identify features, target variables and data quality issues	Dataset understanding and ML problem definition
Phase 3: System Design	Design modular architecture, pipeline workflow, backend and frontend structure	System architecture finalized
Phase 4: Data Analysis & Preprocessing	Implement dataset profiling, missing-value treatment, encoding, scaling and splitting	Automated preprocessing module
Phase 5: ML Development	Implement task detection, model recommendation, model training and evaluation	Automated ML engine
Phase 6: Explainability & Insights	Implement feature importance, natural-language insights and expertise adaptation	Explainable and user-adaptive results
Phase 7: Application Development	Integrate FastAPI backend with Streamlit frontend	Complete interactive application
Phase 8: Testing & Validation	Perform pipeline testing, verify outputs and evaluate model behavior	Validated system
Phase 9: Reporting & Documentation	Generate reports and prepare technical documentation and presentation	Final project and documentation

4.1 Expected Outcomes

1. An automated data profiling mechanism.
2. Automatic machine learning task identification.
3. Automated preprocessing of structured datasets.
4. Intelligent model recommendation.
5. Training and comparison of multiple machine learning algorithms.
6. Automatic selection of the best-performing model.
7. Task-specific model evaluation.
8. Feature-importance-based explainability.
9. Natural-language machine learning insights.
10. Expertise-adaptive explanations.
11. Interactive visualization through Streamlit.
12. API-based access through FastAPI.
13. Persistent storage of trained models.
14. Automatically generated TXT, JSON and PDF reports.

4.2 Future Scope

- Real-time /predict API for inference.
- SHAP-based model-agnostic explainability.
- Hyperparameter optimization using Optuna or GridSearchCV.
- Cross-validation for more robust model evaluation.
- Complete implementation of time-series algorithms.
- Proper training-only fitting of preprocessing transformations.
- Saving the complete preprocessing pipeline along with the trained model.
- MLflow-based experiment and model tracking.
- Docker-based deployment.
- Cloud deployment.
- Database integration for model and experiment history.
- Authentication and security for production APIs.
- Distributed processing using PySpark.
- Streaming industrial telemetry through platforms such as Kafka or Azure Event Hubs.

CONCLUSION

ADSLM – Adaptive Data Science Language Model is an end-to-end automated data science and AutoML system designed to simplify the development and interpretation of machine learning solutions for industrial datasets.

The system integrates dataset analysis, automatic task detection, preprocessing, model recommendation, model training, evaluation, explainability, natural-language insight generation, expertise adaptation, and automated reporting into a unified pipeline.

The predictive-maintenance use case demonstrates how machine sensor information such as temperature, vibration, pressure, RPM, and oil level can be processed to identify machine fault conditions. The combination of machine learning automation and explainable, user-adaptive outputs makes ADSLM useful not only for technical users but also for stakeholders who may not have extensive machine learning knowledge.

The strongest aspect of the project is its adaptability. It adapts to the structure of the dataset, the identified machine learning task, and the expertise level of the user. Consequently, the system goes beyond conventional AutoML by combining automated model development with understandable explanations and structured reporting.

Overall, ADSLM provides a strong foundation for an intelligent industrial data science platform and can be extended in the future toward real-time inference, large-scale data processing, advanced explainability, model lifecycle management, and cloud-based industrial deployment.