

1. Focus sensor calibration and quality assurance on top predictive features: 'job_id'.

2. Investigate extreme sensor spikes detected in columns ['job_id', 'industry_id'] via IQR bounds.
3. R^2 (0.0167) is low; consider engineering non-linear feature interactions.
4. Partition Delta Lake tables on column 'job_id' for sub-second distributed telemetry queries.
5. Deploy the saved model binary via the FastAPI /orchestrate endpoint for production inference.

=====
End of ADSLM Report | Enterprise AI Evaluation
=====